

```
In [1]: %cd C:\Users\manoj\Downloads
```

C:\Users\manoj\Downloads

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [6]: df= pd.read_csv('student_scores.csv')
```

```
In [7]: df.info
```

```
Out[7]: <bound method DataFrame.info of      Hours  Scores
0      2.5      21
1      5.1      47
2      3.2      27
3      8.5      75
4      3.5      30
5      1.5      20
6      9.2      88
7      5.5      60
8      8.3      81
9      2.7      25
10     7.7      85
11     5.9      62
12     4.5      41
13     3.3      42
14     1.1      17
15     8.9      95
16     2.5      30
17     1.9      24
18     6.1      67
19     7.4      69
20     2.7      30
21     4.8      54
22     3.8      35
23     6.9      76
24     7.8      86>
```


```
In [8]: df.head()
```

```
Out[8]:
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [9]: plt.scatter(df['Hours'],df['Scores'])  
plt.xlabel('Hours')  
plt.ylabel('Scores')  
plt.title('Hours vs Scores')
```

```
Out[9]: Text(0.5, 1.0, 'Hours vs Scores')
```



```
In [19]: X=np.array(df.iloc[:, :-1])
```

```
In [24]: y=np.array(df.iloc[:,1])
```

```
In [25]: print(X)
```

```
[[2.5]
 [5.1]
 [3.2]
 [8.5]
 [3.5]
 [1.5]
 [9.2]
 [5.5]
 [8.3]
 [2.7]
 [7.7]
 [5.9]
 [4.5]
 [3.3]
 [1.1]
 [8.9]
 [2.5]
 [1.9]
 [6.1]
 [7.4]
 [2.7]
 [4.8]
 [3.8]
 [6.9]
 [7.8]]
```

```
In [26]: print(y)
```

```
[21 47 27 75 30 20 88 60 81 25 85 62 41 42 17 95 30 24 67 69 30 54 35 76
 86]
```

```
In [27]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
In [28]: from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(X_train, y_train)
```

Out[28]: LinearRegression()

```
In [30]: print(regressor.intercept_)  
  
2.018160041434683
```

```
In [31]: print(regressor.coef_)  
  
[9.91065648]
```

```
In [33]: y_pred=regressor.predict(X_test)
```

```
In [34]: df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})  
df
```

Out[34]:

	Actual	Predicted
0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

```
In [35]: from sklearn import metrics  
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))  
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))  
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

Mean Absolute Error: 4.183859899002975
Mean Squared Error: 21.5987693072174
Root Mean Squared Error: 4.6474476121003665

```
In [36]: df.describe()
```

```
Out[36]:
```

	Actual	Predicted
count	5.000000	5.000000
mean	41.600000	42.651852
std	22.255336	24.407192
min	20.000000	16.884145
25%	27.000000	26.794801
50%	30.000000	33.732261
75%	62.000000	60.491033
max	69.000000	75.357018

```
In [38]: print(y_pred)
```

```
[16.88414476 33.73226078 75.357018    26.79480124 60.49103328]
```

```
In [39]: print(X_test)
```

```
[[1.5]  
 [3.2]  
 [7.4]  
 [2.5]  
 [5.9]]
```

```
In [ ]: #multilinear regression
```

```
In [1]: %cd C:\Users\manoj\Downloads
```

```
C:\Users\manoj\Downloads
```

```
In [11]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [12]: df= pd.read_csv('petrol_consumption.csv')
```

```
In [13]: df.info
```

```
Out[13]: <bound method DataFrame.info of
0      9.00      3571      1976      0.525
1      9.00      4092      1250      0.572
2      9.00      3865      1586      0.580
3      7.50      4870      2351      0.529
4      8.00      4399      431      0.544
5     10.00      5342      1333      0.571
6      8.00      5319     11868      0.451
7      8.00      5126      2138      0.553
8      8.00      4447      8577      0.529
9      7.00      4512      8507      0.552
10     8.00      4391      5939      0.530
11     7.50      5126     14186      0.525
12     7.00      4817      6930      0.574
13     7.00      4207      6580      0.545
14     7.00      4332      8159      0.608
15     7.00      4318     10340      0.586
16     7.00      4206      8508      0.572
17     7.00      3718      4725      0.540
18     7.00      4716      5915      0.724
19     8.50      4341      6010      0.677
20     7.00      4593      7834      0.663
21     8.00      4983       602      0.602
22     9.00      4897      2449      0.511
23     9.00      4258      4686      0.517
24     8.50      4574      2619      0.551
25     9.00      3721      4746      0.544
26     8.00      3448      5399      0.548
27     7.50      3846      9061      0.579
28     8.00      4188      5975      0.563
29     9.00      3601      4650      0.493
30     7.00      3640      6905      0.518
31     7.00      3333      6594      0.513
32     8.00      3063      6524      0.578
33     7.50      3357      4121      0.547
34     8.00      3528      3495      0.487
35     6.58      3802      7834      0.629
36     5.00      4045     17782      0.566
37     7.00      3897      6385      0.586
```

38	8.50	3635	3274	0.663
39	7.00	4345	3905	0.672
40	7.00	4449	4639	0.626
41	7.00	3656	3985	0.563
42	7.00	4300	3635	0.603
43	7.00	3745	2611	0.508
44	6.00	5215	2302	0.672
45	9.00	4476	3942	0.571
46	7.00	4296	4083	0.623
47	7.00	5002	9794	0.593

Petrol_Consumption

0	541
1	524
2	561
3	414
4	410
5	457
6	344
7	467
8	464
9	498
10	580
11	471
12	525
13	508
14	566
15	635
16	603
17	714
18	865
19	640
20	649
21	540
22	464
23	547
24	460
25	566
26	577
27	631
28	574
29	534


```

30          571
31          554
32          577
33          628
34          487
35          644
36          640
37          704
38          648
39          968
40          587
41          699
42          632
43          591
44          782
45          510
46          610
47          524 >

```

In [14]: `df.head()`

Out[14]:

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)	Petrol_Consumption
0	9.0	3571	1976	0.525	541
1	9.0	4092	1250	0.572	524
2	9.0	3865	1586	0.580	561
3	7.5	4870	2351	0.529	414
4	8.0	4399	431	0.544	410

```

In [20]: X = df[['Petrol_tax', 'Average_income', 'Paved_Highways',
                'Population_Driver_licence(%)']]
y = df['Petrol_Consumption']

```

```

In [25]: X=np.array(df.iloc[:,4])
y=np.array(df.iloc[:,4])

```

```
In [26]: print(X)
```

```
[[9.0000e+00 3.5710e+03 1.9760e+03 5.2500e-01]
 [9.0000e+00 4.0920e+03 1.2500e+03 5.7200e-01]
 [9.0000e+00 3.8650e+03 1.5860e+03 5.8000e-01]
 [7.5000e+00 4.8700e+03 2.3510e+03 5.2900e-01]
 [8.0000e+00 4.3990e+03 4.3100e+02 5.4400e-01]
 [1.0000e+01 5.3420e+03 1.3330e+03 5.7100e-01]
 [8.0000e+00 5.3190e+03 1.1868e+04 4.5100e-01]
 [8.0000e+00 5.1260e+03 2.1380e+03 5.5300e-01]
 [8.0000e+00 4.4470e+03 8.5770e+03 5.2900e-01]
 [7.0000e+00 4.5120e+03 8.5070e+03 5.5200e-01]
 [8.0000e+00 4.3910e+03 5.9390e+03 5.3000e-01]
 [7.5000e+00 5.1260e+03 1.4186e+04 5.2500e-01]
 [7.0000e+00 4.8170e+03 6.9300e+03 5.7400e-01]
 [7.0000e+00 4.2070e+03 6.5800e+03 5.4500e-01]
 [7.0000e+00 4.3320e+03 8.1590e+03 6.0800e-01]
 [7.0000e+00 4.3180e+03 1.0340e+04 5.8600e-01]
 [7.0000e+00 4.2060e+03 8.5080e+03 5.7200e-01]
 [7.0000e+00 3.7180e+03 4.7250e+03 5.4000e-01]
 [7.0000e+00 4.7160e+03 5.9150e+03 7.2400e-01]
 [8.5000e+00 4.3410e+03 6.0100e+03 6.7700e-01]
 [7.0000e+00 4.5930e+03 7.8340e+03 6.6300e-01]
 [8.0000e+00 4.9830e+03 6.0200e+02 6.0200e-01]
 [9.0000e+00 4.8970e+03 2.4490e+03 5.1100e-01]
 [9.0000e+00 4.2580e+03 4.6860e+03 5.1700e-01]
 [8.5000e+00 4.5740e+03 2.6190e+03 5.5100e-01]
 [9.0000e+00 3.7210e+03 4.7460e+03 5.4400e-01]
 [8.0000e+00 3.4480e+03 5.3990e+03 5.4800e-01]
 [7.5000e+00 3.8460e+03 9.0610e+03 5.7900e-01]
 [8.0000e+00 4.1880e+03 5.9750e+03 5.6300e-01]
 [9.0000e+00 3.6010e+03 4.6500e+03 4.9300e-01]
 [7.0000e+00 3.6400e+03 6.9050e+03 5.1800e-01]
 [7.0000e+00 3.3330e+03 6.5940e+03 5.1300e-01]
 [8.0000e+00 3.0630e+03 6.5240e+03 5.7800e-01]
 [7.5000e+00 3.3570e+03 4.1210e+03 5.4700e-01]
 [8.0000e+00 3.5280e+03 3.4950e+03 4.8700e-01]
 [6.5800e+00 3.8020e+03 7.8340e+03 6.2900e-01]
 [5.0000e+00 4.0450e+03 1.7782e+04 5.6600e-01]
 [7.0000e+00 3.8970e+03 6.3850e+03 5.8600e-01]
 [8.5000e+00 3.6350e+03 3.2740e+03 6.6300e-01]
```

```
[7.0000e+00 4.3450e+03 3.9050e+03 6.7200e-01]
[7.0000e+00 4.4490e+03 4.6390e+03 6.2600e-01]
[7.0000e+00 3.6560e+03 3.9850e+03 5.6300e-01]
[7.0000e+00 4.3000e+03 3.6350e+03 6.0300e-01]
[7.0000e+00 3.7450e+03 2.6110e+03 5.0800e-01]
[6.0000e+00 5.2150e+03 2.3020e+03 6.7200e-01]
[9.0000e+00 4.4760e+03 3.9420e+03 5.7100e-01]
[7.0000e+00 4.2960e+03 4.0830e+03 6.2300e-01]
[7.0000e+00 5.0020e+03 9.7940e+03 5.9300e-01]]
```

In [27]: `print(y)`

```
[541 524 561 414 410 457 344 467 464 498 580 471 525 508 566 635 603 714
 865 640 649 540 464 547 460 566 577 631 574 534 571 554 577 628 487 644
 640 704 648 968 587 699 632 591 782 510 610 524]
```

In [28]: `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)`

In [32]: `from sklearn.linear_model import LinearRegression`
`regressor = LinearRegression()`
`regressor.fit(X_train, y_train)`

Out[32]: `LinearRegression()`

In [33]: `print(regressor.intercept_)`

```
425.59933220324103
```

In [34]: `print(regressor.coef_)`

```
[-4.00166602e+01 -6.54126674e-02 -4.74073380e-03 1.34186212e+03]
```

In [36]: `y_pred=regressor.predict(X_test)`

```
In [37]: df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
df
```

Out[37]:

	Actual	Predicted
0	534	469.391989
1	410	545.645464
2	577	589.668394
3	571	569.730413
4	577	649.774809
5	704	646.631164
6	487	511.608148
7	587	672.475177
8	467	502.074782
9	580	501.270734

```
In [38]: from sklearn import metrics
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

Mean Absolute Error: 56.822247478964734

Mean Squared Error: 4666.344787588368

Root Mean Squared Error: 68.31064915215173

In [39]: `df.describe()`

Out[39]:

	Actual	Predicted
count	10.000000	10.000000
mean	549.400000	565.827107
std	80.687601	71.811927
min	410.000000	469.391989
25%	498.750000	504.458123
50%	574.000000	557.687939
75%	579.250000	632.390471
max	704.000000	672.475177

In [40]: `print(y_pred)`

```
[469.39198872 545.64546431 589.66839402 569.7304133 649.77480909
646.63116356 511.60814841 672.47517717 502.07478157 501.2707342 ]
```